

<https://www.coursera.org/learn/ai-deep-learning-capstone/ungradedLti/KYMfy/building-a-classifier-with-resnet50>

Exercice lab Building a Classifier with ResNet50

In this lab, you will learn how to build an image classifier using the ResNet50 pre-trained model. You may also download the lab. Click [HERE](#) to download the lab notebook (.ipynb)

Ce cours utilise l'outil d'un tiers, Building a Classifier with ResNet50, pour améliorer votre expérience d'apprentissage. L'outil référence des informations de base comme votre nom, votre adresse e-mail et votre ID Coursera.

- J'accepte d'utiliser cet outil de manière responsable.



Pre-Trained Models

Objective

In this lab, you will learn how to leverage pre-trained models to build image classifiers instead of building a model from scratch.

Table of Contents

1. Import Libraries and Packages
2. Download Data
3. Define Global Constants
4. Construct ImageDataGenerator Instances
5. Compile and Fit Model

Import Libraries and Packages

Let's start the lab by importing the libraries that we will be using in this lab.

First, we will import the ImageDataGenerator module since we will be leveraging it to train our model in batches.

```
[ ]: from keras.preprocessing.image import ImageDataGenerator
    ...
Using TensorFlow backend.
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
    _np_qint8 = np.dtype([("qint8", np.int8, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
    _np_quint8 = np.dtype([("quint8", np.uint8, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:521: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
    _np_qint16 = np.dtype([("qint16", np.int16, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:522: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
    _np_quint16 = np.dtype([("quint16", np.uint16, 1)])
/home/jupyterlab/conda/envs/python/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:523: FutureWarning: Passing (type, 1) or '1type' as a synonym of type is deprecated; in a future ver
```

```
from keras.preprocessing.image import ImageDataGenerator
```

In this lab, we will be using the Keras library to build an image classifier, so let's download the Keras library.

```
[ ]: import keras
from keras.models import Sequential
from keras.layers import Dense
```

```
import keras
from keras.models import Sequential
from keras.layers import Dense
```

Finally, we will be leveraging the ResNet50 model to build our classifier, so let's download it as well.

```
[ ]: from keras.applications import ResNet50
from keras.applications.resnet50 import preprocess_input
    ...
```

```
from keras.applications import ResNet50
from keras.applications.resnet50 import preprocess_input
```

Download Data

For your convenience, I have placed the data on a server which you can retrieve easily using the **wget** command. So let's run the following line of code to get the data. Given the large size of the image dataset, it might take some time depending on your internet speed.

```
[ ]: ## get the data
!wget https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0321EN/data/concrete_data_week3.zip
--2020-11-01 10:11:58--  https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0321EN/data/concrete_data_week3.zip
Resolving s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net)... 6
7.228.254.196
Connecting to s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net)|6
7.228.254.196|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 261482368 (249M) [application/zip]
Saving to: 'concrete_data_week3.zip'

concrete_data_week3 100%[=====] 249.37M 18.0MB/s    in 15s

2020-11-01 10:12:13 (16.7 MB/s) - 'concrete_data_week3.zip' saved [261482368/261482368]

## get the data
!wget https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/DL0321EN/data/concrete\_data\_week3.zip
```

And now if you check the left directory pane, you should see the zipped file *concrete_data_week3.zip* appear. So, let's go ahead and unzip the file to access the images. Given the large number of images in the dataset, this might take a couple of minutes, so please be patient, and wait until the code finishes running.

```
[ ]: !unzip concrete_data_week3.zip
Archive: concrete_data_week3.zip
  creating: concrete_data_week3/
  creating: concrete_data_week3/valid/
  creating: concrete_data_week3/valid/positive/
  inflating: concrete_data_week3/valid/positive/16679_1.jpg
  creating: __MACOSX/concrete_data_week3/
  creating: __MACOSX/concrete_data_week3/valid/
  creating: __MACOSX/concrete_data_week3/valid/positive/
  inflating: __MACOSX/concrete_data_week3/valid/positive/_16679_1.jpg
  inflating: concrete_data_week3/valid/positive/19463.jpg
  inflating: __MACOSX/concrete_data_week3/valid/positive/_19463.jpg
  inflating: concrete_data_week3/valid/positive/19041_1.jpg
```

!unzip concrete_data_week3.zip

Now, you should see the folder *concrete_data_week3* appear in the left pane. If you open this folder by double-clicking on it, you will find that it contains two folders: *train* and *valid*. And if you explore these folders, you will find that each contains two subfolders: *positive* and *negative*. These are the same folders that we saw in the labs in the previous modules of this course, where *negative* is the negative class and it represents the concrete images with no cracks and *positive* is the positive class and it represents the concrete images with cracks.

Important Note: There are thousands and thousands of images in each folder, so please don't attempt to double click on the *negative* and *positive* folders. This may consume all of your memory and you may end up with a **50*** error. So please **DO NOT DO IT**.

Define Global Constants

Here, we will define constants that we will be using throughout the rest of the lab.

1. We are obviously dealing with two classes, so `num_classes` is 2.
2. The ResNet50 model was built and trained using images of size (224 x 224). Therefore, we will have to resize our images from (227 x 227) to (224 x 224).
3. We will training and validating the model using batches of 100 images.

```
[ ]: num_classes = 2  
  
image_resize = 224  
  
batch_size_training = 100  
batch_size_validation = 100
```

`num_classes = 2`

`image_resize = 224`

`batch_size_training = 100`

`batch_size_validation = 100`

Construct ImageDataGenerator Instances

In order to instantiate an `ImageDataGenerator` instance, we will set the `preprocessing_function` argument to `preprocess_input` which we imported from `keras.applications.resnet50` in order to preprocess our images the same way the images used to train ResNet50 model were processed.

```
[ ]: data_generator = ImageDataGenerator(  
    preprocessing_function=preprocess_input,  
)
```

```
data_generator = ImageDataGenerator(  
    preprocessing_function=preprocess_input,  
)
```

Next, we will use the `flow_from_directory` method to get the training images as follows:

```
[ ]: train_generator = data_generator.flow_from_directory(  
    'concrete_data_week3/train',  
    target_size=(image_resize, image_resize),  
    batch_size=batch_size_training,  
    class_mode='categorical')  
Found 30001 images belonging to 2 classes.
```

```
train_generator = data_generator.flow_from_directory(  
    'concrete_data_week3/train',  
    target_size=(image_resize, image_resize),  
    batch_size=batch_size_training,  
    class_mode='categorical')
```

Your Turn: Use the `flow_from_directory` method to get the validation images and assign the result to `validation_generator`.

```
[ ]: ## Type your answer here
```

```
Double-click **here** for the solution.

<!-- The correct answer is:
validation_generator = data_generator.flow_from_directory(
    'concrete_data_week3/valid',
    target_size=(image_resize, image_resize),
    batch_size=batch_size_validation,
    class_mode='categorical')
-->
```

Double-click **here** for the solution.

```
<!-- The correct answer is:
validation_generator = data_generator.flow_from_directory(
    'concrete_data_week3/valid',
    target_size=(image_resize, image_resize),
    batch_size=batch_size_validation,
    class_mode='categorical')
-->
```

Build, Compile and Fit Model

In this section, we will start building our model. We will use the Sequential model class from Keras.

```
[ ]: model = Sequential()
model = Sequential()
```

Next, we will add the ResNet50 pre-trained model to our model. However, note that we don't want to include the top layer or the output layer of the pre-trained model. We actually want to define our own output layer and train it so that it is optimized for our image dataset. In order to leave out the output layer of the pre-trained model, we will use the argument *include_top* and set it to **False**.

```
: model.add(ResNet50(
    include_top=False,
    pooling='avg',
    weights='imagenet',
))
<--- downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.2/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5
94658560/94653016 [=====] - 3s 0us/step
```

```
model.add(ResNet50(
    include_top=False,
    pooling='avg',
    weights='imagenet',
))
```

Then, we will define our output layer as a **Dense** layer, that consists of two nodes and uses the **Softmax** function as the activation function.

```
[ ]: model.add(Dense(num_classes, activation='softmax'))
model.add(Dense(num_classes, activation='softmax'))
```

You can access the model's layers using the *layers* attribute of our model object.

```
[ ]: model.layers  
***  
[12]: [<keras.engine.training.Model at 0x7fed21ac77b8>,  
<keras.layers.core.Dense at 0x7fed80cd8e10>]
```

model.layers

You can see that our model is composed of two sets of layers. The first set is the layers pertaining to ResNet50 and the second set is a single layer, which is our Dense layer that we defined above.

You can access the ResNet50 layers by running the following:

```
[ ]: model.layers[0].layers  
[13]: [<keras.engine.topology.InputLayer at 0x7fed80cd8eb8>,  
<keras.layers.convolutional.ZeroPadding2D at 0x7fed80cf29b0>,  
<keras.layers.convolutional.Conv2D at 0x7fed80cf2e48>,  
<keras.layers.normalization.BatchNormalization at 0x7fed80cf2a20>,  
<keras.layers.core.Activation at 0x7fedaf870f28>,  
<keras.layers.pooling.MaxPooling2D at 0x7fedaef20cf8>,  
<keras.layers.convolutional.Conv2D at 0x7fed80cfc3c8>,  
<keras.layers.normalization.BatchNormalization at 0x7fedaf008400>,  
<keras.layers.core.Activation at 0x7fee4dc46e10>,  
<keras.layers.convolutional.Conv2D at 0x7fee4c9d3438>,  
<keras.layers.normalization.BatchNormalization at 0x7fee4c990c88>,  
<keras.layers.core.Activation at 0x7fee4c9a80f0>.  
<keras.layers.convolutional.Conv2D at 0x7fed21ed3f60>,  
<keras.layers.normalization.BatchNormalization at 0x7fed21ebf6d8>,  
<keras.layers.core.Activation at 0x7fed21e18be0>,  
<keras.layers.convolutional.Conv2D at 0x7fed21dfdeb8>,  
<keras.layers.normalization.BatchNormalization at 0x7fed21deffd0>,  
<keras.layers.core.Activation at 0x7fed21d3f2b0>,  
<keras.layers.convolutional.Conv2D at 0x7fed21c8b748>,  
<keras.layers.normalization.BatchNormalization at 0x7fed21cfb860>,  
<keras.layers.merge.Add at 0x7fed21c2cb38>,  
<keras.layers.core.Activation at 0x7fed21bffd30>,  
<keras.layers.pooling.AveragePooling2D at 0x7fed21bfff28>,
```

model.layers[0].layers

Since the ResNet50 model has already been trained, then we want to tell our model not to bother with training the ResNet part, but to train only our dense output layer. To do that, we run the following.

```
[ ]: model.layers[0].trainable = False
```

model.layers[0].trainable = False

And now using the *summary* attribute of the model, we can see how many parameters we will need to optimize in order to train the output layer.

```
[ ]: model.summary()  
***  
Layer (type)           Output Shape        Param #  
===== ====== =====  
resnet50 (Model)      (None, 2048)       23587712  
dense_1 (Dense)        (None, 2)          4098  
===== ====== =====  
Total params: 23,591,810  
Trainable params: 4,098  
Non-trainable params: 23,587,712
```

```
model.summary()
```

Next we compile our model using the **adam** optimizer.

```
[ ]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

Before we are able to start the training process, with an `ImageDataGenerator`, we will need to define how many steps compose an epoch. Typically, that is the number of images divided by the batch size. Therefore, we define our steps per epoch as follows:

```
[ ]: steps_per_epoch_training = len(train_generator)
steps_per_epoch_validation = len(validation_generator)
num_epochs = 2
```

```
steps_per_epoch_training = len(train_generator)
steps_per_epoch_validation = len(validation_generator)
num_epochs = 2
```

Finally, we are ready to start training our model. Unlike a conventional deep learning training where data is not streamed from a directory, with an `ImageDataGenerator` where data is augmented in batches, we use the **fit_generator** method.

```
[ ]: fit_history = model.fit_generator(
    train_generator,
    steps_per_epoch=steps_per_epoch_training,
    epochs=num_epochs,
    validation_data=validation_generator,
    validation_steps=steps_per_epoch_validation,
    verbose=1,
```

```
fit_history = model.fit_generator(
    train_generator,
    steps_per_epoch=steps_per_epoch_training,
    epochs=num_epochs,
    validation_data=validation_generator,
    validation_steps=steps_per_epoch_validation,
    verbose=1,
)
```

Now that the model is trained, you are ready to start using it to classify images.

Since training can take a long time when building deep learning models, it is always a good idea to save your model once the training is complete if you believe you will be using the model again later. You will be using this model in the next module, so go ahead and save your model.

```
[ ]: model.save('classifier_resnet_model.h5')
```

```
model.save('classifier_resnet_model.h5')
```

Now that the model is trained, you are ready to start using it to classify images.

Since training can take a long time when building deep learning models, it is always a good idea to save your model once the training is complete if you believe you will be using the model again later. You will be using this model in the next module, so go ahead and save your model.

```
[ ]: model.save('classifier_resnet_model.h5')  
model.save('classifier_resnet_model.h5')
```

Now, you should see the model file *classifier_resnet_model.h5* appear in the left directory pane.

Thank you for completing this lab!