

Machine learning: toutes les classifications étudiées

Classification with Python

In this notebook we try to practice all the classification algorithms that we learned in this course.

We load a dataset using Pandas library, and apply the following algorithms, and find the best one for this specific dataset by accuracy evaluation methods.

Lets first load required libraries:

```
[2]: import itertools
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import NullFormatter
import pandas as pd
import numpy as np
import matplotlib.ticker as ticker
from sklearn import preprocessing
%matplotlib inline
```

```
import itertools
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.ticker import NullFormatter
import pandas as pd
import numpy as np
import matplotlib.ticker as ticker
from sklearn import preprocessing
%matplotlib inline
```

About dataset

This dataset is about past loans. The **Loan_train.csv** data set includes details of 346 customers whose loan are already paid off or defaulted. It includes following fields:

Field	Description
Loan_status	Whether a loan is paid off or in collection
Principal	Basic principal loan amount at the
Terms	Origination terms which can be weekly (7 days), biweekly, and monthly payoff schedule
Effective_date	When the loan got originated and took effects
Due_date	Since it's one-time payoff schedule, each loan has one single due date
Age	Age of applicant
Education	Education of applicant
Gender	The gender of applicant

Lets download the dataset

```
[3]: !wget -O loan_train.csv https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/loan_train.csv
--2020-03-29 11:03:14-- https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/loan_train.csv
Resolving s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net)... 67.228.254.196
Connecting to s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net)|67.228.254.196|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 23101 (23K) [text/csv]
Saving to: 'loan_train.csv'

100%[=====] 23,101      --.-K/s   in 0.002s

2020-03-29 11:03:14 (13.3 MB/s) - 'loan_train.csv' saved [23101/23101]
```

```
!wget -O loan_train.csv https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/loan\_train.csv
```

Load Data From CSV File

```
[4]: df = pd.read_csv('loan_train.csv')
df.head()
```

```
[4]:   Unnamed: 0  Unnamed: 0.1  loan_status  Principal  terms  effective_date  due_date  age  education  Gender
0          0          0    PAIDOFF     1000      30  9/8/2016  10/7/2016   45  High School or Below    male
1          2          2    PAIDOFF     1000      30  9/8/2016  10/7/2016   33        Bachelor    female
2          3          3    PAIDOFF     1000      15  9/8/2016  9/22/2016   27         college    male
3          4          4    PAIDOFF     1000      30  9/9/2016  10/8/2016   28         college    female
4          6          6    PAIDOFF     1000      20  9/9/2016  10/8/2016   29         college    male
```

```
df = pd.read_csv('loan_train.csv')
df.head()
```

Convert to date time object

```
[6]: df['due_date'] = pd.to_datetime(df['due_date'])
df['effective_date'] = pd.to_datetime(df['effective_date'])
df.head()
```

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education	Gender
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below	male
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechelor	female
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college	male
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college	female
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college	male

```
df['due_date'] = pd.to_datetime(df['due_date'])
df['effective_date'] = pd.to_datetime(df['effective_date'])
df.head()
```

Data visualization and pre-processing

Let's see how many of each class is in our data set

```
[7]: df['loan_status'].value_counts()
```

loan_status	count
PAIDOFF	260
COLLECTION	86

260 people have paid off the loan on time while 86 have gone into collection

```
df['loan_status'].value_counts()
```

Lets plot some columns to understand data better:

```
[8]: # notice: installing seaborn might takes a few minutes
!conda install -c anaconda seaborn -y
```

Solving environment: done

Package Plan

environment location: /opt/conda/envs/Python36

added / updated specs:

- seaborn

The following packages will be downloaded:

package	build	size	source
ca-certificates-2020.1.1	0	132 KB	anaconda
certifi-2019.11.28	py36_1	157 KB	anaconda
openssl-1.1.1	h7b6447c_0	5.0 MB	anaconda
seaborn-0.10.0	py_0	161 KB	anaconda

```
# notice: installing seaborn might takes a few minutes
!conda install -c anaconda seaborn -y
```

The following packages will be UPDATED:

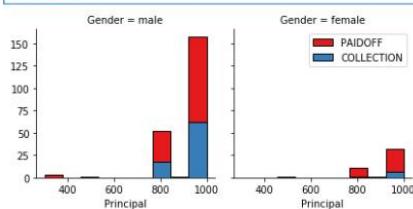
package	old version	new version	source
ca-certificates	2020.1.1-0	2020.1.1-0	anaconda
certifi	2019.11.28-py36_0	2019.11.28-py36_1	anaconda
openssl	1.1.1e-h7b6447c_0	1.1.1-h7b6447c_0	anaconda
seaborn	0.9.0-pyh91ea838_1	0.10.0-py_0	anaconda

Downloading and Extracting Packages
ca-certificates-2020 | 132 KB | ##### | 100%
certifi-2019.11.28 | 157 KB | ##### | 100%
openssl-1.1.1 | 5.0 MB | ##### | 100%
seaborn-0.10.0 | 161 KB | ##### | 100%
Preparing transaction: done
Verifying transaction: done
Executing transaction: done

```
[9]: import seaborn as sns
```

```
bins = np.linspace(df['Principal'].min(), df['Principal'].max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'Principal', bins=bins, ec="k")

g.axes[-1].legend()
plt.show()
```



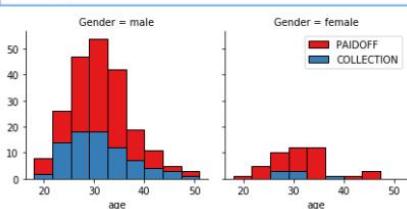
```
import seaborn as sns
```

```
bins = np.linspace(df.Principal.min(), df.Principal.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'Principal', bins=bins, ec="k")

g.axes[-1].legend()
plt.show()
```

```
[10]: bins = np.linspace(df.age.min(), df.age.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'age', bins=bins, ec="k")

g.axes[-1].legend()
```



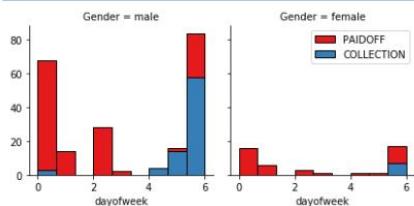
```
bins = np.linspace(df.age.min(), df.age.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'age', bins=bins, ec="k")
```

```
g.axes[-1].legend()  
plt.show()
```

Pre-processing: Feature selection/extraction

Lets look at the day of the week people get the loan

```
[11]: df['dayofweek'] = df['effective_date'].dt.dayofweek
      bins = np.linspace(df['dayofweek'].min(), df['dayofweek'].max(), 10)
      g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
      g.map(plt.hist, 'dayofweek', bins=bins, ec="k")
      g.axes[-1].legend()
      plt.show()
```



```
df['dayofweek'] = df['effective_date'].dt.dayofweek
bins = np.linspace(df.dayofweek.min(), df.dayofweek.max(), 10)
g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
g.map(plt.hist, 'dayofweek', bins=bins, ec="k")
g.axes[-1].legend()
plt.show()
```

We see that people who get the loan at the end of the week dont pay it off, so lets use Feature binarization to set a threshold values less then day 4

```
[12]: df['weekend'] = df['dayofweek'].apply(lambda x: 1 if (x>3) else 0)
df.head()
```

[12]:	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education	Gender	dayofweek	weekend
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below	male	3	0
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechelor	female	3	0
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college	male	3	0
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college	female	4	1
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college	male	4	1

```
df['weekend'] = df['dayofweek'].apply(lambda x: 1 if (x>3) else 0)
df.head()
```

Convert Categorical features to numerical values

Lets look at gender:

```
[13]: df.groupby(['Gender'])['loan_status'].value_counts(normalize=True)
```

```
[13]: Gender loan_status
female PAIDOFF      0.865385
          COLLECTION   0.134615
male   PAIDOFF      0.731293
          COLLECTION   0.268707
Name: loan_status, dtype: float64
```

86 % of female hav there loans while only 73 % of males hav there loan

```
df.groupby(['Gender'])['loan_status'].value_counts(normalize=True)
```

Lets convert male to 0 and female to 1:

```
[14]: df['Gender'].replace(to_replace=['male','female'], value=[0,1],inplace=True)
df.head()
```

```
[14]: Unnamed: 0    Unnamed: 0.1  loan_status  Principal  terms  effective_date  due_date  age  education  Gender  dayofweek  weekend
0            0           0    PAIDOFF     1000     30  2016-09-08  2016-10-07  45  High School or Below    0        3       0
1            2           2    PAIDOFF     1000     30  2016-09-08  2016-10-07  33      Bechelor     1        3       0
2            3           3    PAIDOFF     1000     15  2016-09-08  2016-09-22  27      college     0        3       0
3            4           4    PAIDOFF     1000     30  2016-09-09  2016-10-08  28      college     1        4       1
4            6           6    PAIDOFF     1000     30  2016-09-09  2016-10-08  29      college     0        4       1
```

```
df['Gender'].replace(to_replace=['male','female'], value=[0,1],inplace=True)
df.head()
```

One Hot Encoding

How about education?

```
[15]: df.groupby(['education'])['loan_status'].value_counts(normalize=True)
```

```
[15]: education          loan_status
Bechelor          PAIDOFF      0.750000
                  COLLECTION   0.250000
High School or Below  PAIDOFF      0.741722
                      COLLECTION  0.258278
Master or Above    COLLECTION  0.500000
                  PAIDOFF      0.500000
college           PAIDOFF      0.765101
                      COLLECTION  0.234899
Name: loan_status, dtype: float64
```

```
df.groupby(['education'])['loan_status'].value_counts(normalize=True)
```

Feature before One Hot Encoding

```
[16]: df[['Principal','terms','age','Gender','education']].head()
```

```
[16]: Principal  terms  age  Gender  education
0         1000     30   45      0  High School or Below
1         1000     30   33      1      Bechelor
2         1000     15   27      0      college
3         1000     30   28      1      college
4         1000     30   29      0      college
```

```
df[['Principal','terms','age','Gender','education']].head()
```

Use one hot encoding technique to conver categorical variables to binary variables and append them to the feature Data Frame

```
[17]: Feature = df[['Principal','terms','age','Gender','weekend']]
Feature = pd.concat([Feature,pd.get_dummies(df['education'])], axis=1)
Feature.drop(['Master or Above'], axis = 1,inplace=True)
Feature.head()
```

```
[17]: Principal  terms  age  Gender  weekend  Bechelor  High School or Below  college
0         1000     30   45      0       0       0           1       0
1         1000     30   33      1       0       1           0       0
2         1000     15   27      0       0       0           0       1
3         1000     30   28      1       1       0           0       1
4         1000     30   29      0       1       0           0       1
```

```
Feature = df[['Principal','terms','age','Gender','weekend']]
Feature = pd.concat([Feature,pd.get_dummies(df['education'])], axis=1)
Feature.drop(['Master or Above'], axis = 1,inplace=True)
Feature.head()
```

Feature selection

Lets define feature sets, X:

```
[18]: X = Feature  
X[0:5]
```

	Principal	terms	age	Gender	weekend	Bachelor	High School or Below	college
0	1000	30	45	0	0	0	1	0
1	1000	30	33	1	0	1	0	0
2	1000	15	27	0	0	0	0	1
3	1000	30	28	1	1	0	0	1
4	1000	30	29	0	1	0	0	1

X = Feature

X[0:5]

What are our labels?

```
[20]: y = df['loan_status'].values  
y[0:5]
```

```
[20]: array(['PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF'],  
      dtype=object)
```

y = df['loan_status'].values

y[0:5]

Normalize Data

Data Standardization give data zero mean and unit variance (technically should be done after train test split)

```
[21]: X= preprocessing.StandardScaler().fit(X).transform(X)  
X[0:5]
```

```
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/preprocessing/data.py:645: DataConversionWarning: Data with input dtype uint8, int64 were all converted to float64 by StandardScaler.  
at64 by StandardScaler.  
return self.partial_fit(X, y)  
/opt/conda/envs/Python36/lib/python3.6/site-packages/ipykernel/_main_.py:1: DataConversionWarning: Data with input dtype uint8, int64 were all converted to float64 by StandardScaler.  
if __name__ == '__main__':  
[21]: array([[ 0.51578458,  0.92071769,  2.33152555, -0.42056004, -1.20577805,  
       -0.38170062,  1.13639374, -0.86968108],  
       [ 0.51578458,  0.92071769,  0.34170148,  2.37778177, -1.20577805,  
       2.61985426, -0.87997669, -0.86968108],  
       [ 0.51578458, -0.95911111, -0.65321055, -0.42056004, -1.20577805,  
       -0.38170062, -0.87997669,  1.14984679],  
       [ 0.51578458,  0.92071769, -0.48739188,  2.37778177,  0.82934003,  
       -0.38170062, -0.87997669,  1.14984679],  
       [ 0.51578458,  0.92071769,  0.3215732 , -0.42056004,  0.82934003,  
       -0.38170062, -0.87997669,  1.1498467911])
```

X= preprocessing.StandardScaler().fit(X).transform(X)

X[0:5]

Classification

Now, it is your turn, use the training set to build an accurate model. Then use the test set to report the accuracy of the model. You should use the following algorithm:

- K Nearest Neighbor(KNN)
- Decision Tree
- Support Vector Machine
- Logistic Regression

Notice:

- You can go above and change the pre-processing, feature selection, feature-extraction, and so on, to make a better model.
- You should use either scikit-learn, Scipy or Numpy libraries for developing the classification algorithms.
- You should include the code of the algorithm in the following cells.

K Nearest Neighbor(KNN)

Notice: You should find the best k to build the model with the best accuracy.

warning: You should not use the `loan_test.csv` for finding the best k, however, you can split your `train_loan.csv` into train and test to find the best k.

```
[22]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, random_state=4)  
print ('Train set:', X_train.shape, y_train.shape)  
print ('Test set:', X_test.shape, y_test.shape)
```

```
Train set: (276, 8) (276,)  
Test set: (70, 8) (70,)
```

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2, random_state=4)  
print ('Train set:', X_train.shape, y_train.shape)  
print ('Test set:', X_test.shape, y_test.shape)
```

```
[24]: from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
Ks = 10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))
ConfusionMx = []
for n in range(1,Ks):

    #Train Model and Predict
    neigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat=neigh.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)

    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])

mean_acc
```

```
[24]: array([0.67142857, 0.65714286, 0.71428571, 0.68571429, 0.75714286,
```

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
Ks = 10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))
ConfusionMx = []
for n in range(1,Ks):

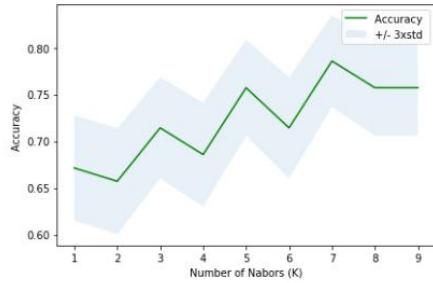
    #Train Model and Predict
    neigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat=neigh.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)

    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])

mean_acc
```

```
[25]: #Plot model accuracy for different number of neighbors
plt.plot(range(1,Ks),mean_acc,'g')
plt.fill_between(range(1,Ks),mean_acc - 1 * std_acc,mean_acc + 1 * std_acc, alpha=0.10)
plt.legend(['Accuracy ', '+/- 3std'])
plt.ylabel('Accuracy ')
plt.xlabel('Number of Nabors (K)')
plt.tight_layout()
plt.show()

print("The best accuracy was with", mean_acc.max(), "with k=", mean_acc.argmax()+1)
```



```
The best accuracy was with 0.7857142857142857 with k= 7
```

```
#Plot model accuracy for different number of neighbors
plt.plot(range(1,Ks),mean_acc,'g')
plt.fill_between(range(1,Ks),mean_acc - 1 * std_acc,mean_acc + 1 * std_acc, alpha=0.10)
plt.legend(['Accuracy ', '+/- 3std'])
plt.ylabel('Accuracy ')
plt.xlabel('Number of Nabors (K)')
plt.tight_layout()
plt.show()

print("The best accuracy was with", mean_acc.max(), "with k=", mean_acc.argmax()+1)
```

Decision Tree

```
[27]: import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
X_trainset, X_testset, y_trainset, y_testset = train_test_split(X, y, test_size=0.3, random_state=3)

[28]: # Modeling
loan_statu_Tree = DecisionTreeClassifier(criterion="entropy", max_depth = 4)
loan_statu_Tree # it shows the default parameters
loan_statu_Tree.fit(X_trainset,y_trainset)

[28]: DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=4,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')
```

27

```
import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
X_trainset, X_testset, y_trainset, y_testset = train_test_split(X, y, test_size=0.3, random_state=3)
```

28

```
# Modeling
loan_statu_Tree = DecisionTreeClassifier(criterion="entropy", max_depth = 4)
loan_statu_Tree # it shows the default parameters
loan_statu_Tree.fit(X_trainset,y_trainset)
```

```
[29]: # Prediction
predTree = loan_statu_Tree.predict(X_testset)
print (predTree [0:5])
print (y_testset [0:5])

['PAIDOFF' 'PAIDOFF' 'PAIDOFF' 'PAIDOFF' 'PAIDOFF']
['PAIDOFF' 'PAIDOFF' 'COLLECTION' 'COLLECTION' 'PAIDOFF']
```

```
# Prediction
predTree = loan_statu_Tree.predict(X_testset)
print (predTree [0:5])
print (y_testset [0:5])
```

```
[30]: # Evaluation
from sklearn import metrics
import matplotlib.pyplot as plt
print("DecisionTrees's Accuracy: ", metrics.accuracy_score(y_testset, predTree))

DecisionTrees's Accuracy: 0.6538461538461530
```

```
# Evaluation
from sklearn import metrics
import matplotlib.pyplot as plt
print("DecisionTrees's Accuracy: ", metrics.accuracy_score(y_testset, predTree))
```

```
[35]: # Visualization
!conda install -c conda-forge pydotplus -y
!conda install -c conda-forge python-graphviz -y

from sklearn.externals.six import StringIO
import pydotplus
import matplotlib.image as mpimg
from sklearn import tree
%matplotlib inline

Solving environment: done
```

```
# Visualization
!conda install -c conda-forge pydotplus -y
!conda install -c conda-forge python-graphviz -y
```

```
from sklearn.externals.six import StringIO
import pydotplus
import matplotlib.image as mpimg
from sklearn import tree
```

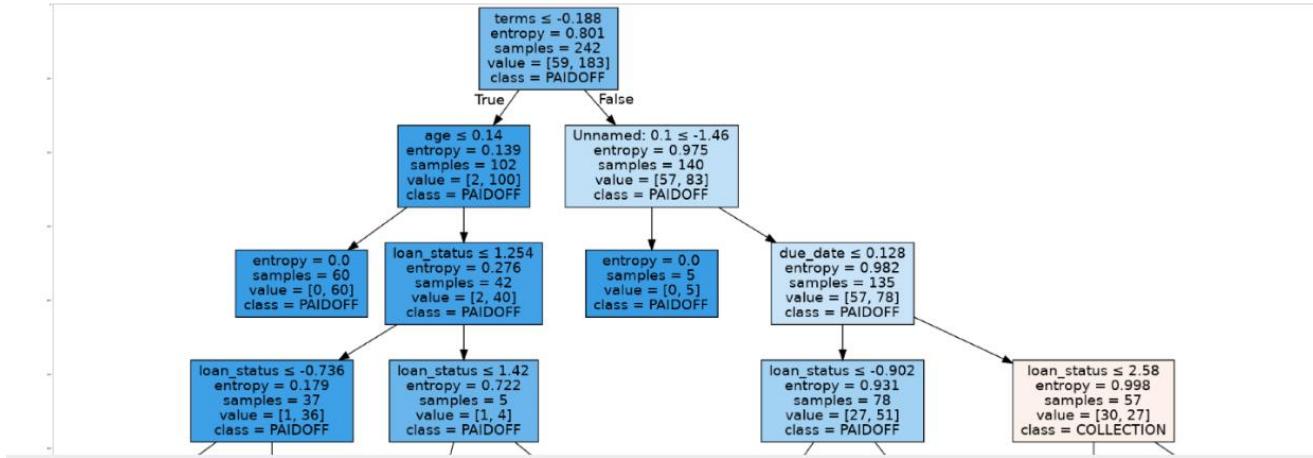
```
%matplotlib inline
```

```
# All requested packages already installed.
```

```
Solving environment: done
```

```
# All requested packages already installed.
```

```
[4]: dot_data = StringIO()
filename = "loan_statu_Tree.png"
featureNames = df.columns[0:8]
targetNames = df['loan_status'].unique().tolist()
out=tree.export_graphviz(loan_statu_Tree,feature_names=featureNames, out_file=dot_data, class_names= np.unique(y_trainset), filled=True, special_characters=True,rotate=False)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png(filename)
img = mpimg.imread(filename)
plt.figure(figsize=(100, 200))
plt.imshow(img,interpolation='nearest')
```



```
dot_data = StringIO()
filename = "loan_statu_Tree.png"
featureNames = df.columns[0:8]
targetNames = df['loan_status'].unique().tolist()
out=tree.export_graphviz(loan_statu_Tree,feature_names=featureNames, out_file=dot_data, class_names= np.unique(y_trainset), filled=True, special_characters=True,rotate=False)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png(filename)
img = mpimg.imread(filename)
plt.figure(figsize=(100, 200))
plt.imshow(img,interpolation='nearest')
```

Support Vector Machine

```
[37]: import pandas as pd
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
%matplotlib inline
import matplotlib.pyplot as plt

# Change df['Loan_Status'] values to 0 or 1 int
df['loan_status'].replace(to_replace=['PAIDOFF','COLLECTION'], value=[1,0],inplace=True)
df.head()
```

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education	Gender	dayofweek	weekend
0	0	0	1	1000	30	2016-09-08	2016-10-07	45	High School or Below	0	3	0
1	2	2	1	1000	30	2016-09-08	2016-10-07	33	Bachelor	1	3	0
2	3	3	1	1000	15	2016-09-08	2016-09-22	27	college	0	3	0
3	4	4	1	1000	30	2016-09-09	2016-10-08	28	college	1	4	1
4	6	6	1	1000	30	2016-09-09	2016-10-08	29	college	0	4	1

```
import pandas as pd
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
%matplotlib inline
```

```

import matplotlib.pyplot as plt

# Change df['loan_status'] values to 0 or 1 int
df['loan_status'].replace(to_replace=['PAIDOFF','COLLECTION'], value=[1,0],inplace=True)
df.head()

```

```
[40]: # Concatenate feature with new values of df['loan_status']
New_Feature = pd.concat([Feature,df['loan_status']], axis=1)
```

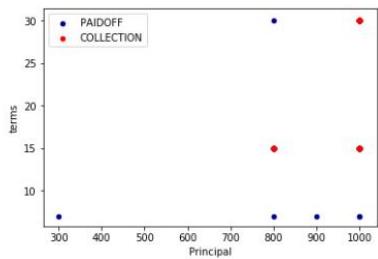
```
New_Feature.head()
```

	Principal	terms	age	Gender	weekend	Bachelor	High School or Below	college	loan_status
0	1000	30	45	0	0	0	1	0	1
1	1000	30	33	1	0	1	0	0	1
2	1000	15	27	0	0	0	0	1	1
3	1000	30	28	1	1	0	0	1	1
4	1000	30	29	0	1	0	0	1	1

```
# Concatenate feature with new values of df['loan_status']
New_Feature = pd.concat([Feature,df['loan_status']], axis=1)
```

```
New_Feature.head()
```

```
[42]: # Let's look at the distribution
ax = New_Feature[New_Feature['loan_status'] == 1][0:50].plot(kind='scatter', x='Principal', y='terms', color='DarkBlue', label='PAIDOFF');
New_Feature[New_Feature['loan_status'] == 0][0:50].plot(kind='scatter', x='Principal', y='terms', color='Red', label='COLLECTION', ax=ax);
plt.show()
```



```
# Let's look at the distribution
ax = New_Feature[New_Feature['loan_status'] == 1][0:50].plot(kind='scatter', x='Principal', y='terms', color='DarkBlue',
label='PAIDOFF');
New_Feature[New_Feature['loan_status'] == 0][0:50].plot(kind='scatter', x='Principal', y='terms', color='Red',
label='COLLECTION', ax=ax);
plt.show()
```

```
[43]: New_Feature.dtypes
```

Principal		int64
terms		int64
age		int64
Gender		int64
weekend		int64
Bachelor		uint8
High School or Below		uint8
college		uint8
loan_status		int64
dtype:	object	

```
New_Feature.dtypes
```

```
[44]: # Drop the Non numerical rows in Bechelor, College and High school columns
New_Feature = New_Feature[pd.to_numeric(New_Feature['Bechelor'], errors='coerce').notnull()]
New_Feature['Bechelor'] = New_Feature['Bechelor'].astype('int')
New_Feature = New_Feature[pd.to_numeric(New_Feature['High School or Below'], errors='coerce').notnull()]
New_Feature['High School or Below'] = New_Feature['High School or Below'].astype('int')
New_Feature = New_Feature[pd.to_numeric(New_Feature['college'], errors='coerce').notnull()]
New_Feature['college'] = New_Feature['college'].astype('int')

New_Feature.dtypes
```

```
[44]: Principal          int64
terms            int64
age              int64
Gender           int64
weekend          int64
Bechelor          int64
High School or Below    int64
college          int64
loan_status       int64
dtype: object
```

```
# Drop the Non numerical rows in Bechelor, College and High school columns
New_Feature = New_Feature[pd.to_numeric(New_Feature['Bechelor'], errors='coerce').notnull()]
New_Feature['Bechelor'] = New_Feature['Bechelor'].astype('int')
New_Feature = New_Feature[pd.to_numeric(New_Feature['High School or Below'], errors='coerce').notnull()]
New_Feature['High School or Below'] = New_Feature['High School or Below'].astype('int')
New_Feature = New_Feature[pd.to_numeric(New_Feature['college'], errors='coerce').notnull()]
New_Feature['college'] = New_Feature['college'].astype('int')
```

New_Feature.dtypes

```
[45]: # X Data selection
feature2_df = New_Feature[['Principal', 'terms', 'age', 'Gender', 'weekend', 'Bechelor', 'High School or Below', 'college']]
X1 = np.asarray(feature2_df)
X1[0:5]
```

```
[45]: array([[1000,   30,   45,   0,   0,   0,   1,   0],
       [1000,   30,   33,   1,   0,   1,   0,   0],
       [1000,   15,   27,   0,   0,   0,   0,   1],
       [1000,   30,   28,   1,   1,   0,   0,   1],
       [1000,   30,   29,   0,   1,   0,   0,   1]])
```

X Data selection

```
feature2_df = New_Feature[['Principal', 'terms', 'age', 'Gender', 'weekend', 'Bechelor', 'High School or Below', 'college']]
X1 = np.asarray(feature2_df)
X1[0:5]
```

```
[50]: # y Data selection
New_Feature['loan_status'] = New_Feature['loan_status'].astype('int')
y1 = np.asarray(New_Feature['loan_status'])
y1 [0:5]
```

```
[50]: array([1, 1, 1, 1, 1])
```

y Data selection

```
New_Feature['loan_status'] = New_Feature['loan_status'].astype('int')
y1 = np.asarray(New_Feature['loan_status'])
y1 [0:5]
```

```
[51]: # Train and test data
X1_train, X1_test, y1_train, y1_test = train_test_split( X1, y1, test_size=0.2, random_state=4)
print ('Train set:', X1_train.shape, y1_train.shape)
print ('Test set:', X1_test.shape, y1_test.shape)
```

```
Train set: (276, 8) (276,)
Test set: (70, 8) (70,)
```

Train and test data

```
X1_train, X1_test, y1_train, y1_test = train_test_split( X1, y1, test_size=0.2, random_state=4)
print ('Train set:', X1_train.shape, y1_train.shape)
print ('Test set:', X1_test.shape, y1_test.shape)
```

```
[52]: # Modeling
from sklearn import svm
clf = svm.SVC(kernel='rbf')
clf.fit(X1_train, y1_train)
yhat = clf.predict(X1_test)
yhat [0:5]

/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/svm/base.py:196: FutureWarning: The default value of gamma will change from 'auto' to 'scale' in version 0.22 to
account better for unscaled features. Set gamma explicitly to 'auto' or 'scale' to avoid this warning.
"avoid this warning.", FutureWarning)
```

```
# Modeling
from sklearn import svm
clf = svm.SVC(kernel='rbf')
clf.fit(X1_train, y1_train)
yhat = clf.predict(X1_test)
yhat [0:5]
```

```
[53]: # Evaluation
from sklearn.metrics import classification_report, confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                         normalize=False,
                         title='Confusion matrix',
                         cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
```

```
# Evaluation
from sklearn.metrics import classification_report, confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                         normalize=False,
                         title='Confusion matrix',
                         cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
```

```

thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

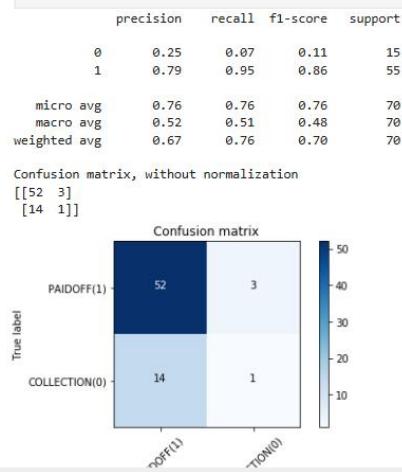
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

```
[55]: # Compute the confusion matrix
cnf_matrix = confusion_matrix(y1_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)

print (classification_report(y1_test, yhat))

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['PAIDOFF(1)','COLLECTION(0)'],normalize= False, title='Confusion matrix')
```



```

# Compute the confusion matrix
cnf_matrix = confusion_matrix(y1_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)

print (classification_report(y1_test, yhat))

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['PAIDOFF(1)','COLLECTION(0)'],normalize= False, title='Confusion matrix')

```

Logistic Regression

```
[57]: # Normalize the dataset
from sklearn import preprocessing
X1_norm = preprocessing.StandardScaler().fit(X1).transform(X1)
X1_norm[0:5]
```

/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/utils/validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/utils/validation.py:595: DataConversionWarning: Data with input dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)

```
# Normalize the dataset
from sklearn import preprocessing
X1_norm = preprocessing.StandardScaler().fit(X1).transform(X1)
X1_norm[0:5]
```

```
[58]: # Train and test dataset
from sklearn.model_selection import train_test_split
X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_state=4)
print('Train set:', X1_train.shape, y1_train.shape)
print('Test set:', X1_test.shape, y1_test.shape)
```

Train set: (276, 8) (276,)
Test set: (70, 8) (70,)

Train and test dataset

```
from sklearn.model_selection import train_test_split
X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_state=4)
print('Train set:', X1_train.shape, y1_train.shape)
print('Test set:', X1_test.shape, y1_test.shape)
```

```
[59]: # Modeling Logistic regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
LR = LogisticRegression(C=0.01, solver='liblinear')
LR
```

```
[59]: LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='liblinear',
    tol=0.0001, verbose=0, warm_start=False)
```

Modeling logistic regression

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import confusion_matrix
```

```
LR = LogisticRegression(C=0.01, solver='liblinear').fit(X1_train,y1_train)
```

LR

```
[60]: # Predicting using our test set  
y1hat = LR.predict(X1_test)  
y1hat
```

```
# Predicting using our test set
```

```
y1hat = LR.predict(X1_test)
```

y1hat

```
[61]: # Predict proba estimate all classes ordered by label  
y1hat_prob = LR.predict_proba(X1_test)
```

```
[61]: array([[0.32, 0.68],  
           [0.25, 0.75],  
           [0.18, 0.82],  
           [0.32, 0.68],  
           [0.24, 0.76],  
           [0.21, 0.79],  
           [0.27, 0.73],  
           [0.24, 0.76],  
           [0.32, 0.68],  
           [0.3 , 0.7 ],  
           [0.3 , 0.7 ],  
           [0.32, 0.68],  
           [0.29, 0.71],  
           [0.3 , 0.7 ],  
           [0.14, 0.86],  
           [0.16, 0.84],  
           [0.4 , 0.6 ],  
           [0.18, 0.82],  
           [0.34, 0.66],  
           [0.22, 0.78],  
           [0.28, 0.72],  
           [0.31, 0.69],  
           [0.36, 0.64],  
           [0.37, 0.63],  
           [0.21, 0.79],  
           [0.34, 0.66],  
           [0.35, 0.65],  
           [0.15, 0.85],  
           [0.35, 0.65],  
           [0.14, 0.86],  
           [0.22, 0.78],  
           [0.34, 0.66],  
           [0.28, 0.72],  
           [0.28, 0.72],  
           [0.21, 0.79],  
           [0.2 , 0.8 ],  
           [0.34, 0.66],  
           [0.14, 0.86],  
           [0.31, 0.69],  
           [0.21, 0.79],  
           [0.35, 0.65],  
           [0.23, 0.77],  
           [0.17, 0.83],  
           [0.34, 0.66],  
           [0.19, 0.81],  
           [0.34, 0.66],  
           [0.23, 0.77],  
           [0.32, 0.68],  
           [0.34, 0.66],  
           [0.19, 0.81],  
           [0.34, 0.66],  
           [0.23, 0.77],  
           [0.24, 0.76],  
           [0.2 , 0.8 ],  
           [0.24, 0.76],  
           [0.29, 0.71],  
           [0.26, 0.74],  
           [0.1 , 0.9 ],  
           [0.21, 0.79],  
           [0.25, 0.75],  
           [0.22, 0.78],  
           [0.27, 0.73],  
           [0.12, 0.88],  
           [0.28, 0.72],  
           [0.2 , 0.8 ],  
           [0.41, 0.59],  
           [0.24, 0.76],  
           [0.31, 0.69],  
           [0.28, 0.72],  
           [0.2 , 0.8 ],  
           [0.16, 0.84],  
           [0.24, 0.76]]))
```

```
# Predict proba estimate all classes ordered by label
```

```
y1hat_prob = LR.predict_proba(X1_test)  
y1hat_prob
```

Model Evaluation using Test set

```
[62]: from sklearn.metrics import jaccard_similarity_score  
from sklearn.metrics import f1_score  
from sklearn.metrics import log_loss
```

First, download and load the test set:

```
from sklearn.metrics import jaccard_similarity_score  
from sklearn.metrics import f1_score  
from sklearn.metrics import log_loss
```

```
[63]: !wget -O loan_test.csv https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/loan_test.csv
--2020-03-29 16:33:58-- https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/loan_test.csv
Resolving s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net)... 67.228.254.196
Connecting to s3-api.us-geo.objectstorage.softlayer.net (s3-api.us-geo.objectstorage.softlayer.net)|67.228.254.196|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 3642 (3.6K) [text/csv]
Saving to: 'loan_test.csv'

100%[=====] 3,642      --.-K/s   in 0s

2020-03-29 16:33:58 (335 MB/s) - 'loan_test.csv' saved [3642/3642]
```

!wget -O loan_test.csv https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/CognitiveClass/ML0101ENv3/labs/loan_test.csv

Load Test set for evaluation

```
[64]: test_df = pd.read_csv('loan_test.csv')
test_df.head()
```

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education	Gender
0	1	1	PAIDOFF	1000	30	9/8/2016	10/7/2016	50	Bechelor	female
1	5	5	PAIDOFF	300	7	9/9/2016	9/15/2016	35	Master or Above	male
2	21	21	PAIDOFF	1000	30	9/10/2016	10/9/2016	43	High School or Below	female
3	24	24	PAIDOFF	1000	30	9/10/2016	10/9/2016	26	college	male
4	35	35	PAIDOFF	800	15	9/11/2016	9/25/2016	29	Bechelor	male

```
test_df = pd.read_csv('loan_test.csv')
test_df.head()
```

```
[65]: # Evaluate the model using jaccard
from sklearn.metrics import jaccard_similarity_score
jaccard_similarity_score(y1_test, y1hat)
```

```
[65]: 0.7857142857142857
```

```
# Evaluate the model using jaccard
from sklearn.metrics import jaccard_similarity_score
jaccard_similarity_score(y1_test, y1hat)
```

```
[66]: # Evaluating the confusing matrix
from sklearn.metrics import classification_report, confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                        normalize=False,
                        title='Confusion matrix',
                        cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

    cnf_matrix = confusion_matrix(y1_test, yhat, labels=[1,0])
    np.set_printoptions(precision=2)

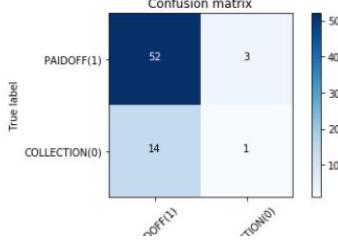
    print (classification_report(y1_test, yhat))

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['PAIDOFF(1)', 'COLLECTION(0)'], normalize= False, title='Confusion matrix')
```

	precision	recall	f1-score	support
0	0.25	0.07	0.11	15
1	0.79	0.95	0.86	55

	micro avg	macro avg	weighted avg	
0	0.76	0.52	0.67	70
1	0.76	0.51	0.70	70

Confusion matrix, without normalization
[[52 3]
 [14 1]]



```
# Evaluating the confusing matrix
from sklearn.metrics import classification_report, confusion_matrix
import itertools
```

```
def plot_confusion_matrix(cm, classes,
                        normalize=False,
                        title='Confusion matrix',
                        cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
```

```

plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

cnf_matrix = confusion_matrix(y1_test, yhat, labels=[1,0])
np.set_printoptions(precision=2)

print(classification_report(y1_test, yhat))

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=['PAIDOFF(1)','COLLECTION(0)'],normalize= False, title='Confusion matrix')

```

```
[69]: # Evaluating the model using log_loss
log_loss(y1_test, y1hat_prob)
```

```
[69]: 0.5435595463662589
```

```
# Evaluating the model using log_loss
log_loss(y1_test, y1hat_prob)
```

```
[71]: # 1. Using jaccard to evaluate KNN
from sklearn.metrics import jaccard_similarity_score
jaccard_similarity_score(y_test, yhat)
```

```
[71]: 0.7571428571428571
```

```
# 1. Using jaccard to evaluate KNN
from sklearn.metrics import jaccard_similarity_score
jaccard_similarity_score(y_test, yhat)
```

```
[75]: # Using jaccard to evaluate Decision tree
jaccard_similarity_score(y_testset, predTree)
```

```
[75]: 0.6538461538461539
```

```
# Using jaccard to evaluate Decision tree
jaccard_similarity_score(y_testset, predTree)
```

```
[76]: # Using jaccard to evaluate SVM
jaccard_similarity_score(y1_test, y1hat)
```

```
[76]: 0.7857142857142857
```

```
# Using jaccard to evaluate SVM
jaccard_similarity_score(y1_test, y1hat)
```

```
[77]: # Using jaccard to evaluate Logistic regression
jaccard_similarity_score(y1_test, y1hat)
```

```
[77]: 0.7857142857142857
```

```
# Using jaccard to evaluate Logistic regression
jaccard_similarity_score(y1_test, y1hat)
```

```
[81]: # Using jaccard to evaluate different models
print('jaccard to accurate KNN =',jaccard_similarity_score(y_test, yhat))
print('jaccard to accurate Decision Tree =',jaccard_similarity_score(y_testset, predTree))
print('jaccard to accurate SVM =',jaccard_similarity_score(y1_test, y1hat))
print('jaccard to accurate Logistic Regression =',jaccard_similarity_score(y1_test, y1hat))
```

```
jaccard to accurate KNN = 0.7571428571428571
jaccard to accurate Decision Tree = 0.6538461538461539
jaccard to accurate SVM = 0.7857142857142857
jaccard to accurate Logistic Regression = 0.7857142857142857
```

```
# Using jaccard to evaluate different models
print('jaccard to accurate KNN =',jaccard_similarity_score(y_test, yhat))
print('jaccard to accurate Decision Tree =',jaccard_similarity_score(y_testset, predTree))
print('jaccard to accurate SVM =',jaccard_similarity_score(y1_test, y1hat))
print('jaccard to accurate Logistic Regression =',jaccard_similarity_score(y1_test, y1hat))
```

```
[88]: # Using f1 score to evaluate different models
print('F1 score to accurate KNN')
print (classification_report(y_test, yhat))
print('F1 score to accurate Decision Tree')
print (classification_report(y_testset, predTree))
print('F1 score to accurate SVM')
print (classification_report(y1_test, y1hat))
print('F1 score to accurate Logistic Regression')
print (classification_report(y1_test, y1hat))

F1 score to accurate KNN
precision    recall  f1-score   support
          0       0.25      0.07      0.11      15
          1       0.79      0.95      0.86      55
micro avg       0.76      0.76      0.76      70
macro avg       0.52      0.51      0.48      70
weighted avg    0.67      0.76      0.70      70

F1 score to accurate Decision Tree
precision    recall  f1-score   support
COLLECTION    0.37      0.48      0.42      27
PAIDOFF       0.80      0.71      0.75      77
micro avg       0.65      0.65      0.65     104
weighted avg    0.69      0.65      0.67     104

F1 score to accurate SVM
precision    recall  f1-score   support
          0       0.00      0.00      0.00      15
          1       0.79      1.00      0.88      55
micro avg       0.79      0.79      0.79      70
macro avg       0.39      0.50      0.44      70
weighted avg    0.62      0.79      0.69      70

F1 score to accurate Logistic Regression
precision    recall  f1-score   support
          0       0.00      0.00      0.00      15
          1       0.79      1.00      0.88      55
micro avg       0.79      0.79      0.79      70
macro avg       0.39      0.50      0.44      70
weighted avg    0.62      0.79      0.69      70

/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)
/opt/conda/envs/Python36/lib/python3.6/site-packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples.
  'precision', 'predicted', average, warn_for)
```

```
# Using f1 score to evaluate different models
print('F1 score to accurate KNN')
print (classification_report(y_test, yhat))
print('F1 score to accurate Decision Tree')
print (classification_report(y_testset, predTree))
print('F1 score to accurate SVM')
print (classification_report(y1_test, y1hat))
print('F1 score to accurate Logistic Regression')
print (classification_report(y1_test, y1hat))
```

Report

You should be able to report the accuracy of the built model using different evaluation metrics:

Algorithm	Jaccard	F1-score	LogLoss
KNN	?	?	NA
Decision Tree	?	?	NA
SVM	?	?	NA
LogisticRegression	?	?	?

```
[105]: # Create a report table
import pandas as pd
df = pd.DataFrame({'Jaccard': [0.7571428571428571, 0.6538461538461539, 0.7857142857142857, 0.7857142857142857], \
                   'F1-score': [0.70, 0.67 , 0.69, 0.69], \
                   'LogLoss': ['NA', 'NA', 'NA',0.5435595463662589]}, \
                   index = ['KNN','Decision Tree','SVM','LogisticRegression'])
df.head()
```

```
[105]:      Jaccard  F1-score  LogLoss
KNN    0.757143    0.70     NA
Decision Tree  0.653846    0.67     NA
SVM    0.785714    0.69     NA
LogisticRegression  0.785714    0.69   0.54356
```

```
# Create a report table
import pandas as pd
df = pd.DataFrame({'Jaccard': [0.7571428571428571, 0.6538461538461539,
0.7857142857142857,0.7857142857142857], \
                   'F1-score': [0.70, 0.67 , 0.69, 0.69], \
                   'LogLoss': ['NA', 'NA', 'NA',0.5435595463662589]}, \
                   index = ['KNN','Decision Tree','SVM','LogisticRegression'])
df.head()
```

Want to learn more?

IBM SPSS Modeler is a comprehensive analytics platform that has many machine learning algorithms. It has been designed to bring predictive intelligence to decisions made by individuals, by groups, by systems – by your enterprise as a whole. A free trial is available through this course, available here: [SPSS Modeler](#)

Also, you can use Watson Studio to run these notebooks faster with bigger datasets. Watson Studio is IBM's leading cloud solution for data scientists, built by data scientists. With Jupyter notebooks, RStudio, Apache Spark and popular libraries pre-packaged in the cloud, Watson Studio enables data scientists to collaborate on their projects without having to install anything. Join the fast-growing community of Watson Studio users today with a free account at [Watson Studio](#)

Thanks for completing this lesson!

Author: Saeed Aghabozorgi

Saeed Aghabozorgi, PhD is a Data Scientist in IBM with a track record of developing enterprise level applications that substantially increases clients' ability to turn data into actionable knowledge. He is a researcher in data mining field and expert in developing advanced analytic methods like machine learning and statistical modelling on large datasets.

Copyright © 2018 Cognitive Class. This notebook and its source code are released under the terms of the MIT License.